


# **ASSESSMENT**

## **ON :**

# **TERRAFORM**


## **2**

**TO  
THE  
NEW**



## 1. Launch an ASG in AWS and do Rolling Deployment with change in User Data in LaunchConfig using terraform.

STEP 1: User-data (empty now)

A terminal window with a dark purple background. The prompt is `#!/bin/bash` in cyan. Below it are three tilde characters `~` on separate lines, also in cyan.

STEP 2: main.tf

```
variable "aws_region" {}
variable "template_name" {}
variable "ami" {}
variable "instance_type" {}
variable "key" {}
variable "az" {}
variable "sg" {}
variable "auto_scaling_policy_up" {}
variable "auto_scaling_policy_down" {}
variable "asg" {}
provider "aws" {
    region = var.aws_region
}
resource "aws_launch_template" "template" {
    name = var.template_name
    image_id = var.ami
    instance_initiated_shutdown_behavior = "terminate"
    instance_type = var.instance_type
    key_name = var.key
    placement {
        availability_zone = var.az
    }
    vpc_security_group_ids = [var.sg]
    user_data = filebase64("userdata.sh")
    lifecycle {
        create_before_destroy = true
    }
}
resource "aws_autoscaling_group" "asg" {
    name = var.asg
    availability_zones = ["us-east-1a"]
}
```

```

        desired_capacity = 1
        max_size = 2
        min_size = 1
        health_check_grace_period = 300
        health_check_type = "EC2"
        force_delete = true
        launch_template {
            id = aws_launch_template.template.id
            version = "$Latest"
        }
        vpc_zone_identifier = ["subnet-7857a027"]
    }

    resource "aws_autoscaling_policy" "up_policy" {
        name = var.auto_scaling_policy_up
        scaling_adjustment = 1
        adjustment_type = "ChangeInCapacity"
        cooldown = 300
        autoscaling_group_name = aws_autoscaling_group.asg.name
    }

    resource "aws_autoscaling_policy" "down_policy" {
        name = var.auto_scaling_policy_down
        scaling_adjustment = -1
        adjustment_type = "ChangeInCapacity"
        cooldown = 300
        autoscaling_group_name = aws_autoscaling_group.asg.name
    }

    resource "aws_cloudwatch_metric_alarm" "cpu-high" {
        alarm_name = "High-CPU-Utilization"
        comparison_operator = "GreaterThanOrEqualToThreshold"
        evaluation_periods = 2
        metric_name = "CPU-Utilization"
        period = 300
        statistic = "Average"
        threshold = 80
        # Namespace must never be NULL
        namespace = "System/Linux"
        alarm_actions = [aws_autoscaling_policy.up_policy.arn]
        dimensions = {
            AutoScalingGroupName = aws_autoscaling_group.asg.name
        }
    }
}

```

### STEP 3: env.tfvars

```
aws_region = "us-east-1"
template_name = "Test-Template"
ami = "ami-0020db9a2596c5437"
instance_type = "t2.micro"
key = "diksha_aws"
az = "us-east-1a"
sg = "sg-bbbfa695"
auto_scaling_policy_up = "Up-Scaling-Policy"
auto_scaling_policy_down = "Down-Scaling-Policy"
asg = "Test-ASG"
```

### STEP 4: terraform plan and apply

```
diksha@diksha:~/terraform2$ terraform plan --var-file='env.tfvars'
```

### STEP 5: Check console for ASG and all the infrastructure

Filter: <input type="text" value="Filter Auto Scaling groups..."/>										1 to 1 of	
<input checked="" type="checkbox"/>	Name	Launch Configuration /	Instances	Desired	Min	Max	Availability Zones	Default Cooldown	Health Check Grace		
<input checked="" type="checkbox"/>	Test-ASG	Test-Template	1	1	1	2	us-east-1a	300	300		

\*\*\*\*\*

## 2. Deploy a sample nginx/tomcat/react service on it.

### STEP 1: Make changes in the user-data

```
#!/bin/bash
sudo apt-get update -y
sudo apt-get install nginx -y
systemctl restart nginx
```

### STEP 2: Running # terraform apply again:

-> One change = Launch Template Modified

```
Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

  Enter a value: yes

aws_launch_template.template: Modifying... [id=lt-08b6d35fec10eee5d]
aws_launch_template.template: Modifications complete after 4s [id=lt-08b6d35fec10eee5d]

Apply complete! Resources: 0 added, 1 changed, 0 destroyed.
diksha@diksha:~/terraform2$
```

### STEP 3: Checking User Data:



**View/Change User Data**

Instance ID: i-034264949dce3d318

User Data:

```
#!/bin/bash
apt-get update -y
apt-get install nginx -y
systemctl restart nginx
```

3.86.149.185

Maps (3) TO THE NE... (3) Diksha To... Learning | Das... Linux Academy

## Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](https://nginx.org).  
Commercial support is available at [nginx.com](https://nginx.com).

*Thank you for using nginx.*

\*\*\*\*\*

**3. Attach a LB and create R53 endpoint pointing to lab, service should be accessible from the endpoint.**

**STEP 1:** Creating Application Load Balancer, Target Group and Listener Rules and creating Private Hosted Zone and Record Set:

```
resource "aws_lb" "lb" {
  name = var.lb_name
  internal = "false"
  load_balancer_type = "application"
  enable_cross_zone_load_balancing = true
  security_groups = [var.sg]
  subnets = ["subnet-00986421", "subnet-0981bb37"]
}
resource "aws_lb_target_group" "tg" {
```

```

        name = var.target_group
        port = 80
        protocol = "HTTP"
        vpc_id = "vpc-306e6c4a"
    }
    resource "aws_lb_listener" "listener" {
        load_balancer_arn = aws_lb.lb.arn
        port = 80
        protocol = "HTTP"
        default_action {
            type = "forward"
            target_group_arn = aws_lb_target_group.tg.arn
        }
    }
    # Attatching Target-Group with Auto-Scaling Group ( Here, the Target of
    # Target-Group is Auto-Scaling-Group and not Instances alone )
    resource "aws_autoscaling_attachment" "attach_tg" {
        autoscaling_group_name = aws_autoscaling_group.asg.name
        alb_target_group_arn = aws_lb_target_group.tg.arn
    }

    resource "aws_route53_zone" "private_zone" {
        name = var.private_zone
        vpc {
            vpc_id = "vpc-306e6c4a"
        }
    }
    resource "aws_route53_record" "www" {
        zone_id = aws_route53_zone.private_zone.id
        name = var.record_name
        type = "A"
        alias {
            name = aws_lb.lb.dns_name
            zone_id = aws_lb.lb.zone_id
            evaluate_target_health = false
        }
    }
}

```

**STEP 2:** Add the below variables in env.tfvars

```

lb_name = "Test-Alb"
target_group = "Test-Target-Group"
private_zone = "diksha.com"
record_name = "www.diksha.com"

```

### STEP 3: Apply terraform

```
aws_route53_record.www: Still creating... [10s elapsed]
aws_route53_record.www: Still creating... [20s elapsed]
aws_route53_record.www: Still creating... [30s elapsed]
aws_route53_record.www: Still creating... [40s elapsed]
aws_route53_record.www: Still creating... [50s elapsed]
aws_route53_record.www: Still creating... [1m0s elapsed]
aws_route53_record.www: Still creating... [1m10s elapsed]
aws_route53_record.www: Creation complete after 1m20s [id=Z03913292JP365WZM
6V5K_www.diksha.com_A]

Apply complete! Resources: 6 added, 0 changed, 0 destroyed.
```

### STEP 4: Check for the DNS that you created

Create Hosted Zone

Go to Record Sets

Delete Hosted Zone

Search all fields

X

All Types

Domain Name

Type

Record Set Count

Comment

diksha.com.

Private

3

Managed by Terraform

### STEP 5: Checking the Output using DNS:

```
ubuntu@ip-172-31-95-233:~$ curl www.diksha.com
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

#### 4. Variablize all parameters and pass values as env.tfvars file.

```
diksha@diksha:~/terraform2$ cat env.tfvars
aws_region = "us-east-1"
template_name = "Test-Template"
ami = "ami-07ebfd5b3428b6f4d"
instance_type = "t2.micro"
key = "diksha_aws"
az = "us-east-1a" sg = "sg-bbbfa695"
auto_scaling_policy_up = "Up-Scaling-Policy"
auto_scaling_policy_down = "Down-Scaling-Policy"
asg = "Test-ASG"
lb_name = "Test-Alb"
target_group = "Test-Target-Group"
private_zone = "diksha.com"
record_name = "www.diksha.com"
```

In main.tf specify all variables at the top

```
diksha@diksha:~/terraform2$ cat main.tf variable "aws_region" {}
variable "template_name" {}
variable "ami" {}
variable "instance_type" {}
variable "key" {}
variable "az" {}
variable "sg" {}
variable "auto_scaling_policy_up" {}
variable "auto_scaling_policy_down" {}
variable "asg" {}
variable "lb_name" {}
variable "target_group" {}
variable "private_zone" {}
variable "record_name" {}
```



**5. Create ASG from Launch Template and use a mix of on demand and on spot instance type in the ASG. Instance Type for On Demand and Spot should be different.**

**STEP 1:** Destroyed previous Infrastructure>Editing main.tf ( Changing the aws\_autoscaling\_group configuration):

The Instance Types of On-Demand and Spot-Instances are also Different.

```
#Requesting for Spot Instance of "type = c4.large" only
resource "aws_spot_instance_request" "cheap" {
  ami = var.ami
  spot_price = "0.03"
  instance_type = "c4.xlarge"
```

**6.Enable Spot Feature to use multiple instance type if requested instance type is not available.**

```
# If we want some other Spot-Instances if our Required Instance type is not
available, then we can use Fleet-Request as below
resource "aws_spot_fleet_request" "cheap_compute" {
  iam_fleet_role = var.fleet_role
  spot_price = "0.03"
  allocation_strategy = "diversified"
  target_capacity = 1
  valid_until = "2019-11-04T20:44:20Z"
  launch_specification {
    instance_type = "m4.10xlarge"
    ami = var.ami
    spot_price = "2.793"
    key_name = var.key
    availability_zone = var.az
    subnet_id = "subnet-7857a027"
    weighted_capacity = 1
  }
  launch_specification {
    instance_type = "m4.4xlarge"
    ami = var.ami
    key_name = var.key
    spot_price = "1.117"
    availability_zone = var.az
    subnet_id = "subnet-7857a027"
    weighted_capacity = 1
  }
}
```

```

#Requesting for Spot Instance of "type = c4.large" only
resource "aws_spot_instance_request" "cheap" {
    ami = var.ami
    spot_price = "0.03"
    instance_type = "c4.xlarge"
}

# If we want some other Spot-Instances if our Required Instance type i
s not available, then we can use Fleet-Request as below
resource "aws_spot_fleet_request" "cheap_compute" {
    iam_fleet_role = var.fleet_role
    spot_price = "0.03"
    allocation_strategy = "diversified"
    target_capacity = 1
    valid_until = "2019-11-04T20:44:20Z"
    launch_specification {
        instance_type = "m4.10xlarge"
        ami = var.ami
        spot_price = "2.793"
        key_name = var.key
        availability_zone = var.az
        subnet_id = "subnet-7857a027"
        weighted_capacity = 1
    }
    launch_specification {
        instance_type = "m4.4xlarge"
        ami = var.ami
        key_name = var.key
        spot_price = "1.117"
        availability_zone = var.az
        subnet_id = "subnet-7857a027"
        weighted_capacity = 1
    }
}

```